

Multi-Model Ensemble Methods from Climate Science: A Technical Review for Computational Materials Science Applications

1. Introduction and Cross-Domain Structural Parallels

1.1 The GLIM-CMIP6 Analogy

The structural isomorphism between climate modeling and computational materials science represents one of the most promising opportunities for methodological cross-pollination in modern scientific computing. The **GLIM (Generalized Lattice Information for Materials) platform**, which systematically benchmarks **23 interatomic potentials across 12,000 materials for elastic constant predictions**, embodies an ensemble architecture that is fundamentally identical to how **CMIP6 (Coupled Model Intercomparison Project Phase 6)** compares approximately **30 climate models across global spatial grid points** ([Semantic Scholar](#)) . This parallel extends far beyond superficial similarities in numbers to encompass the core challenges of combining multiple imperfect models, each with distinct structural assumptions, parameterizations, and domains of applicability, into unified predictions with well-calibrated uncertainty estimates.

The **ensemble architecture** in both domains comprises multiple independent modeling approaches applied to common prediction targets. In climate science, these targets are spatiotemporal fields of temperature, precipitation, pressure, and other meteorological variables; in materials science, they are scalar or tensor material properties such as **elastic constants, formation energies, and phonon spectra**. The dimensionality differs—climate models operate on millions of grid points across time, while GLIM evaluates thousands of discrete materials—but the **statistical structure of the ensemble prediction problem remains analogous**. Both domains require methods that can handle high-dimensional prediction spaces where model performance varies systematically, and where this variation itself contains actionable information about model strengths and weaknesses ([Academia.edu](#)) .

The **spatial benchmarking dimension** strengthens this analogy considerably. CMIP6 evaluates models against observational data at each grid point, creating a spatial field of performance metrics. Similarly, GLIM evaluates interatomic potentials against reference data (typically density functional theory calculations or experimental measurements) for each material in its benchmark set. This creates a **materials-property field of performance metrics** that can be analyzed using spatial statistics methods developed for climate model evaluation. The critical insight is that both domains involve high-dimensional prediction spaces where **model performance varies systematically across the prediction domain**, and where this variation must inform ensemble combination and uncertainty quantification ([Academia.edu](#)) .

Feature	Climate Science (CMIP6)	Materials Science (GLIM)
Ensemble size	~30 climate models	23 interatomic potentials
Prediction targets	Temperature, precipitation, pressure fields	Elastic constants (tensor/scalar)
Spatial domain	~64,000 global grid points	12,000 materials
Reference data	Instrumental observations, reanalysis	DFT calculations, experiments
Core challenge	Structural uncertainty, model dependence	Same

Table 1 – continued

Feature	Climate Science (CMIP6)	Materials Science (GLIM)
Temporal structure	Sequential forecasts, lead-time dependent	Static predictions, cross-sectional

Table 1: Structural parallels between CMIP6 climate ensembles and GLIM materials ensembles.

1.2 Core Challenges in Both Domains

1.2.1 Model Diversity vs. Redundancy in Ensemble Construction A persistent tension in both fields arises from the need to balance **model diversity** (capturing structural uncertainty through genuinely different approaches) against **model redundancy** (avoiding pseudoreplication that inflates apparent confidence without adding information). Climate models share numerical methods, parameterization schemes, and even substantial code bases due to historical development patterns and limited modeling groups. The CMIP6 ensemble contains **37 uniquely named models but only 19 independent model families** when structural relationships are accounted for (GMD) . Similarly, interatomic potentials cluster by functional form: multiple **EAM (Embedded Atom Method)** parameterizations, **SNAP (Spectral Neighbor Analysis Potential)** variants with similar bispectrum descriptors, and **neural network potentials** trained on overlapping DFT databases may appear distinct while producing highly correlated errors (eLib - DLR electronic library) .

The consequences of unacknowledged redundancy are **systematically overconfident uncertainty estimates**. When 15 of 23 GLIM potentials share similar functional forms or training data, their agreement provides far less confirmation than 15 truly independent predictions would, yet naive ensemble averaging treats all agreements equally. The **effective sample size** of a dependent ensemble can be much smaller than its nominal size, with uncertainty estimates that are too narrow by factors of two or more (Academia.edu) .

1.2.2 The “Model Democracy” Problem and Equal-Weight Averaging Limitations The “**model democracy**” approach—assigning equal weight to all ensemble members regardless of demonstrated performance—has been the default in both fields, yet is now recognized as suboptimal. In climate science, this practice assumes that models are (a) reasonably independent, (b) equally plausible, (c) distributed around reality, and (d) that their spread represents true uncertainty. As **Knutti and Tebaldi (2010)** demonstrated, **none of these conditions holds strictly**: models share code and ideas, vary dramatically in skill, exhibit common structural biases, and may collectively miss key processes (eLib - DLR electronic library) . The same critique applies to interatomic potential ensembles, where equal weighting ignores vast differences in accuracy across material classes and fails to account for shared training data or methodological heritage (Academia.edu) .

The limitations of equal-weight averaging are particularly severe when **ensemble performance is heterogeneous**. If one machine learning potential consistently outperforms empirical potentials for transition metal oxides, while the reverse holds for simple metals, global equal weighting degrades performance in both regimes. **Performance-based weighting** that adapts to local (material-specific) skill offers substantial improvements, but requires careful implementation to avoid overfitting and to properly account for model dependence (Semantic Scholar) .

1.2.3 Temporal vs. Compositional Generalization Challenges A fundamental asymmetry between the domains lies in the nature of generalization required. **Climate models must generalize**

across time: they are validated on historical observations and applied to future projections, requiring extrapolation in forcing space (greenhouse gas concentrations, aerosol emissions) where the system may exhibit nonlinear responses. **Interatomic potentials must generalize across compositional and structural space:** they are trained and validated on specific material classes and applied to predict properties of novel compositions or structures, requiring extrapolation in chemical space where bonding environments may differ qualitatively from training data ([Academia.edu](#)) .

Despite this difference, the **mathematical challenges are analogous:** detecting when predictions enter **out-of-distribution regimes**, quantifying how uncertainty grows with extrapolation distance, and ensuring that ensemble spread appropriately signals prediction difficulty. Both domains benefit from **ensemble methods that can flag untested domains** and inflate uncertainty accordingly, rather than providing overconfident predictions based on interpolation performance.

2. Bayesian Model Averaging (BMA) for Weighted Ensemble Combination

2.1 Theoretical Foundation

2.1.1 Posterior Model Probability Framework Bayesian Model Averaging (BMA) provides a principled statistical framework for combining ensemble predictions that explicitly accounts for **model uncertainty** through probabilistic weighting. Unlike ad hoc weighting schemes or simple averaging, BMA treats model selection as an inference problem, where the probability of each model given the data determines its contribution to the final prediction. This approach is particularly valuable when multiple models show comparable but imperfect performance, and when the goal is robust prediction rather than model identification ([Semantic Scholar](#)) .

The foundational insight of BMA, developed by **Raftery, Gneiting, and colleagues**, is that **conditioning on a single “best” model ignores uncertainty about model structure**, leading to overconfident predictions and underestimated uncertainty. By averaging over the model space, BMA incorporates this **structural uncertainty** into predictive distributions, producing wider, more realistic prediction intervals. This property is especially critical in materials science, where the true interatomic potential is never known exactly, and any finite ensemble represents an incomplete sampling of possible functional forms ([Semantic Scholar](#)) .

2.1.2 Predictive Distribution as Weighted Mixture of Member Distributions In the BMA framework, the predictive distribution is a **finite mixture** where each ensemble member contributes its predictive distribution weighted by its posterior probability. This mixture structure has profound implications for uncertainty quantification: the total predictive variance decomposes into **within-model variance** (average uncertainty of individual models) and **between-model variance** (disagreement among models), ensuring that model disagreement contributes appropriately to overall uncertainty. This decomposition is unavailable from any single model or from simple point-averaging methods ([Semantic Scholar](#)) .

2.1.3 BMA Weights as Posterior Probabilities Incorporating Model Uncertainty The BMA weights are not arbitrary coefficients but **genuine posterior probabilities** that balance goodness-of-fit against model complexity through the marginal likelihood. Models that achieve better predictive performance without excessive complexity receive higher weights, implementing an **automatic Occam’s razor**. This automatic trade-off is particularly valuable when ensemble members vary substantially in their effective number of parameters, as is common when comparing classical empirical potentials with highly flexible machine learning potentials ([Semantic Scholar](#)) .

2.2 Mathematical Formulation

2.2.1 Predictive Density: $p(y|D) = \sum_k p(y|M_k, D) \cdot p(M_k|D)$ The core BMA equation expresses the predictive distribution for a quantity of interest y given data D as:

$$p(y|D) = \sum_{k=1}^K p(y|M_k, D) \cdot p(M_k|D)$$

where $p(y|M_k, D)$ is the predictive distribution under model M_k and $p(M_k|D)$ is the posterior probability of model M_k (Semantic Scholar) . This formulation ensures that BMA predictions are **never worse than single-model predictions in expectation**, with equality only when one model has posterior probability approaching unity.

The posterior model probabilities are computed via **Bayes' theorem**:

$$p(M_k|D) = \frac{p(D|M_k) \cdot p(M_k)}{\sum_{l=1}^K p(D|M_l) \cdot p(M_l)}$$

where $p(D|M_k) = \int p(D|\theta_k, M_k) p(\theta_k|M_k) d\theta_k$ is the **marginal likelihood** (integrated over model parameters) and $p(M_k)$ is the prior model probability, typically uniform when no strong prior information exists (Semantic Scholar) .

2.2.2 Member-Specific Weight Estimation via EM Algorithm or MCMC In practice, computing marginal likelihoods is analytically intractable for complex models. **Raftery et al. (2005)** developed an efficient **Expectation-Maximization (EM) algorithm** for estimating BMA weights by maximum likelihood, treating model indicators as latent variables (Semantic Scholar) . The algorithm iterates between:

- **E-step:** Compute posterior probabilities of model membership given current parameters
- **M-step:** Update weights and parameters using weighted maximum likelihood

For large ensembles or complex posterior distributions, **Markov Chain Monte Carlo (MCMC)** methods provide an alternative for exploring the full posterior distribution over model weights, enabling uncertainty quantification for the weights themselves (Semantic Scholar) .

2.2.3 Variance Decomposition: Within-Model and Between-Model Components For **Gaussian predictive distributions** with member-specific means μ_k and variances σ_k^2 , the BMA predictive distribution has:

$$\mu_{BMA} = \sum_{k=1}^K w_k \mu_k$$
$$\sigma_{BMA}^2 = \underbrace{\sum_{k=1}^K w_k \sigma_k^2}_{\text{within-model variance}} + \underbrace{\sum_{k=1}^K w_k (\mu_k - \mu_{BMA})^2}_{\text{between-model variance}}$$

where $w_k = p(M_k|D)$. This **explicit variance decomposition** is critical for uncertainty quantification: it separates irreducible uncertainty from model disagreement, guiding targeted investments in model improvement (Semantic Scholar) .

2.3 Climate Science Applications and Extensions

2.3.1 Handling Exchangeable and Missing Ensemble Members A critical extension developed by **Fraley, Raftery, and Gneiting (2010)** addresses **exchangeable and missing ensemble members**—common complications in operational forecasting and CMIP archives ([Semantic Scholar](#)) . **Exchangeability** arises when subsets of ensemble members are statistically indistinguishable, such as multiple runs from the same modeling center or perturbations around a single model initialization. The extended framework treats exchangeable members as drawn from a common predictive distribution, with **group-level weights** that pool information and reduce effective parameter count.

Missing data is ubiquitous in CMIP: not all models provide output for all variables, scenarios, or time periods. The BMA extension uses **imputation strategies** that marginalize over missing values or use available subset information for partial likelihood evaluation, maximizing use of expensive computed data ([Semantic Scholar](#)) .

2.3.2 Stratified BMA for Precipitation Forecasting Standard BMA with Gaussian components is inappropriate for **precipitation**, which is nonnegative, often zero, and right-skewed. **Slughter et al. (2007)** developed **stratified BMA** using **gamma-Gaussian mixture models**: a point mass at zero for dry events and a gamma distribution for positive precipitation amounts ([ResearchGate](#)) . This **regime-specific adaptation**—different statistical treatment for different prediction domains—parallels the materials science challenge of treating different bonding types or structure classes with appropriately tailored models.

2.3.3 Performance in CMIP5/CMIP6 Multi-Model Ensembles BMA applications to CMIP5/CMIP6 demonstrate **consistent improvements** over equal-weight averaging. For temperature projections, BMA reduces mean squared error by **10-20%** compared to the multi-model mean, with larger improvements for precipitation where model disagreement is greater. Perhaps more importantly, BMA produces **well-calibrated probabilistic predictions**: prediction intervals achieve their nominal coverage, meaning 90% credible intervals contain the true value approximately 90% of the time ([Semantic Scholar](#)) .

However, challenges arise for **long-range climate projections** due to **non-stationarity**: historical performance may imperfectly predict future skill, and the “truth” against which to calibrate is itself evolving. This has motivated research into **time-varying BMA weights** and robustness checks for stationarity assumptions ([OUCI](#)) .

2.4 Transfer Potential to Materials Science

2.4.1 Weighting Interatomic Potentials by Predictive Accuracy The direct application of BMA to GLIM-scale interatomic potential ensembles involves treating each potential as a “model” and elastic constant predictions as “forecasts.” The **training data** for weight estimation comprises benchmark materials with high-confidence reference values (typically **DFT with validated exchange-correlation functionals**). Potentials with smaller and more consistent errors receive higher marginal likelihoods and thus higher BMA weights, with **automatic downweighting of poorly performing potentials** ([Semantic Scholar](#)) .

A key advantage is **adaptive weighting by material class**: BMA can estimate separate weights for metals, oxides, semiconductors, etc., if performance varies systematically. This is implemented through **stratified estimation** or by including material descriptors as covariates in the likelihood specification.

2.4.2 Addressing Missing Data in Incomplete Potential Benchmarks Missing data is common in potential benchmarking: **not all potentials support all elements or crystal structures**. The **Fraley et al. (2010) extension** for exchangeable and missing members is directly applicable, using available subset information for partial likelihood evaluation. For materials where only 15 of 23 potentials provide predictions, BMA weights are estimated from this incomplete set, with appropriate uncertainty inflation to account for reduced information ([Semantic Scholar](#)) .

2.4.3 Computational Tractability for Large-Scale Materials Databases The **EM algorithm for BMA weight estimation scales linearly** with ensemble size and dataset size, making it feasible for GLIM-scale applications (23 potentials × 12,000 materials = 276,000 predictions). Once weights are estimated, prediction for new materials requires only **weighted averaging of precomputed potential outputs**—negligible cost compared to potential evaluation itself. This efficiency enables **real-time uncertainty quantification** for high-throughput screening workflows ([Semantic Scholar](#)) .

Aspect	Climate Application	Materials Application
Ensemble members	Climate models (CMIP6)	Interatomic potentials (GLIM)
Training data	Historical observations	DFT benchmark calculations
Target variable	Temperature, precipitation	Elastic constants (Cij, K, G)
Missing data	Variable/scenario gaps	Element/structure incompatibilities
Stratification	Season, region, lead time	Material class, crystal system
Key advantage	Calibrated uncertainties	Adaptive weighting by chemistry

Table 2: BMA implementation comparison between climate and materials science.

3. Reliability Diagrams and Probabilistic Calibration

3.1 Reliability Assessment Framework

3.1.1 Definition: Observed Frequency vs. Forecast Probability Plots Reliability diagrams constitute a fundamental diagnostic for assessing whether **probabilistic forecasts are statistically consistent with observed outcomes**. The construction is conceptually straightforward: for a set of probabilistic forecasts, partition the forecast probability range into discrete bins, compute the average forecast probability and observed event frequency within each bin, and plot these against each other. **Perfect reliability** corresponds to the one-to-one diagonal: events forecast with probability p occur with frequency p ([AGU Publications](#)) .

3.1.2 Perfect Reliability and the One-to-One Reference Line The **one-to-one reference line** in a reliability diagram represents the ideal of **probabilistic calibration**. Deviations from this line reveal systematic miscalibration patterns that guide post-processing interventions. The visual simplicity of reliability diagrams belies their diagnostic power: experienced practitioners can diagnose specific model deficiencies from characteristic curve shapes ([AGU Publications](#)) .

3.1.3 Underconfidence and Overconfidence Detection Systematic deviations from the diagonal indicate specific calibration problems:

Pattern	Interpretation	Typical Cause
Below diagonal	Overconfidence: predicted probabilities too extreme	Ensemble spread too small, model overfitting
Above diagonal	Underconfidence: predicted probabilities too moderate	Ensemble spread too large, excessive regularization
S-shaped curve	Poor discrimination at probability extremes	Inadequate model resolution, threshold effects
Flat curve	No resolution: forecasts near climatology	Overly conservative ensemble, missing predictability

Table 3: Reliability diagram patterns and their interpretation.

3.2 Mathematical Construction

3.2.1 Probability Binning and Frequency Estimation For **binary events**, reliability diagram construction proceeds as follows. Given N forecast-observation pairs with predicted probabilities $\{p_i\}_{i=1}^N$ and binary outcomes $\{o_i\}_{i=1}^N$:

1. Define B bins with edges $0 = b_0 < b_1 < \dots < b_B = 1$
2. Assign each forecast to bin b if $b_{b-1} \leq p_i < b_b$
3. Compute **mean predicted probability** in each bin: $\bar{p}_b = \frac{1}{N_b} \sum_{i \in \text{bin } b} p_i$
4. Compute **observed frequency** in each bin: $\bar{o}_b = \frac{1}{N_b} \sum_{i \in \text{bin } b} o_i$

The reliability diagram plots $(\bar{p}_b, \bar{o}_b)$ points, with consistency bars indicating sampling uncertainty ([AGU Publications](#)) .

3.2.2 Brier Score Decomposition: Reliability - Resolution + Uncertainty The **Brier score** (mean squared probability error) decomposes into three interpretable components ([AGU Publications](#)) :

$$\text{BS} = \underbrace{\sum_{b=1}^B \frac{N_b}{N} (\bar{p}_b - \bar{o}_b)^2}_{\text{Reliability}} - \underbrace{\sum_{b=1}^B \frac{N_b}{N} (\bar{o}_b - \bar{o})^2}_{\text{Resolution}} + \underbrace{\bar{o}(1 - \bar{o})}_{\text{Uncertainty}}$$

- **Reliability:** Measures calibration quality (lower is better; zero for perfect calibration)
- **Resolution:** Measures discrimination between different probability regimes (higher is better)
- **Uncertainty:** Represents irreducible variance of the climatological distribution

This decomposition guides improvement efforts: poor reliability can be addressed through recalibration, while poor resolution indicates fundamental limitations in predictive information content ([AGU Publications](#)) .

3.2.3 Poisson-Binomial Distribution for Exact Reliability Computation For **small samples or extreme probabilities**, the standard binning approach can be improved using the **Poisson-Binomial distribution**, which describes the distribution of sums of independent but non-identically distributed Bernoulli trials. Unlike the Binomial distribution, which assumes constant probability, the Poisson-Binomial accounts for **heterogeneity of forecast probabilities**, providing more accurate confidence intervals and hypothesis tests for reliability ([AGU Publications](#)) .

The **CORP (Consistent, Optimally binned, Reproducible, and Pool-adjacent-violators) approach** represents state-of-the-art reliability diagram construction. Based on **nonparametric isotonic**

regression via the **Pool-Adjacent-Violators (PAV) algorithm**, CORP generates reliability diagrams that are provably consistent and optimally binned without ad hoc choices ([arXiv.org](#)) .

3.3 Calibration Methods

3.3.1 Parametric Recalibration via Logistic Regression **Logistic recalibration** fits a calibration function mapping raw forecast probabilities to calibrated probabilities:

$$\text{logit}(p_{\text{cal}}) = a + b \cdot \text{logit}(p_{\text{raw}})$$

The **slope parameter** b corrects reliability: $b < 1$ for overconfident forecasts (spread reduction), $b > 1$ for underconfident forecasts (spread inflation). The **intercept** a corrects systematic bias. Parameters are estimated by maximum likelihood on training data ([AGU Publications](#)) .

3.3.2 Non-Parametric Isotonic Regression Approaches **Isotonic regression** finds the **monotonic function** that best fits reliability diagram points in a least-squares sense, with no parametric assumptions beyond monotonicity. The **PAV algorithm** efficiently computes this solution, providing flexible recalibration that can capture complex reliability patterns. Isotonic regression typically outperforms parametric methods when sufficient training data are available, but requires larger samples to avoid overfitting ([arXiv.org](#)) .

3.3.3 Integration with BMA and EMOS for Calibrated Probabilistic Outputs Rather than **post-hoc recalibration** of ensemble probabilities, **BMA and EMOS produce calibrated predictions by construction**—provided their model assumptions are satisfied. The reliability diagram then serves as a **diagnostic for model misspecification** rather than a basis for recalibration. This integrated approach, standard in operational weather forecasting, ensures that probabilistic outputs are ready for decision-making without additional processing ([Semantic Scholar](#)) .

3.4 Materials Science Adaptation

3.4.1 Defining “Observed” Elastic Constants for Calibration A critical adaptation challenge is defining **appropriate reference data** for reliability assessment. Unlike weather forecasting with direct instrumental measurements, materials “observations” are typically **DFT calculations** with their own systematic errors relative to experiment. A **hierarchical approach** treats high-level DFT (e.g., PBEsol with converged parameters) as the calibration target, while acknowledging this as an approximation to true physical properties. **Multi-fidelity extensions** incorporate multiple reference sources with appropriate uncertainty attribution ([Academia.edu](#)) .

3.4.2 Binning Strategies for Continuous Property Predictions For **continuous elastic constants**, reliability assessment requires **discretization of the prediction space**. Effective strategies include:

- **Value-based binning:** Partition by predicted bulk modulus ranges (0-50, 50-100, 100-150 GPa, etc.)
- **Quantile-based binning:** Equal-frequency bins ensuring adequate sample sizes
- **Material-class stratification:** Separate reliability diagrams for metals, oxides, semiconductors
- **Confidence-based binning:** Assess whether high-confidence predictions are appropriately reliable

The resulting “**reliability surfaces**” reveal how calibration varies across property space, identifying domains where specific potentials are trustworthy ([AGU Publications](#)) .

3.4.3 Reliability-Driven Potential Selection and Combination Reliability information enables **adaptive ensemble construction**: potentials demonstrating poor reliability in specific regimes can be **excluded or downweighted** for those regimes. More sophisticated approaches use reliability to **select the most appropriate single potential** for each prediction, or to **combine potentials with regime-dependent weights**. This adaptive combination can outperform static weighting by responding to varying applicability domains ([Academia.edu](#)) .

4. Model Independence Testing and Quantification

4.1 The Independence Problem in Ensembles

4.1.1 Pseudoreplication from Shared Model Heritage The assumption of **model independence** is systematically violated in practice, with **profound consequences for uncertainty quantification**. **Pseudoreplication**—the illusion of independent information from dependent sources—arises from **shared code heritage, common training data, overlapping development teams, and shared theoretical foundations**. In climate science, models branch from common ancestors, share parameterization schemes, and are tuned to similar targets. In materials science, machine learning potentials trained on overlapping DFT databases may appear distinct while producing **highly correlated errors** ([Academia.edu](#)) .

4.1.2 Consequences of Dependence: Overconfident Uncertainty Estimates When **K models share substantial common errors**, their agreement provides **far less than K-fold confirmation**, yet standard analyses treat all agreements equally. The **effective sample size** of a dependent ensemble can be **much smaller than its nominal size**, with uncertainty estimates that are **too narrow by factors of two or more**. This overconfidence is particularly insidious because it is not visually apparent: an ensemble of 30 correlated models produces smooth, apparently robust predictions with misleadingly narrow confidence intervals ([Academia.edu](#)) .

4.1.3 Distinguishing Structural from Parametric Model Differences **Structural differences** (alternative functional forms, different fundamental approximations) may produce more independent errors than **parametric differences** (variations within a common structure). Two neural network potentials with different architectures but shared training data may be more dependent than two classical potentials with different functional forms. **Independence should be assessed in prediction space**, not in model formulation space: models with similar error patterns are effectively dependent regardless of surface differences ([eLib - DLR electronic library](#)) .

4.2 The ClimWIP Framework

4.2.1 Performance-and-Independence Weighting Scheme The **Climate model Weighting by Independence and Performance (ClimWIP)** framework, developed by **Knutti and colleagues**, provides systematic integration of **performance-based and independence-based weighting** ([esm-valtool.org](#)) . The method addresses the fundamental limitation of performance-only weighting: rewarding accuracy without penalizing redundancy can concentrate weight on clusters of similar models, failing to capture structural uncertainty.

4.2.2 Inter-Model Distance Metrics ClimWIP defines **pairwise distance in prediction space** as the foundation for independence quantification. For climate models, the **spread distance** is:

$$S_{ij} = \sqrt{(\text{SAT}\Delta_i - \text{SAT}\Delta_j)^2 + (\text{PR}\Delta_i - \text{PR}\Delta_j)^2}$$

where $\text{SAT}\Delta_i$ and $\text{PR}\Delta_i$ represent **root-mean-square errors** of model i against observations for surface air temperature and precipitation patterns (esmvaltool.org). Models with small S_{ij} make similar predictions and are judged less independent.

The **intermember distance metric** I_{ij} normalizes by ensemble spread:

$$I_{ij} = \frac{1}{2} \sum_{l=1}^2 \left(\frac{\phi_{ij}}{\bar{\phi}} \right)_l$$

where ϕ_{ij} is the RMSE between models i and j for predictor field l , and $\bar{\phi}$ is the ensemble mean RMSE (esmvaltool.org).

4.2.3 Normalized Independence Score For a subset of models $s_1, \dots, s_n$, the **normalized independence score** aggregates pairwise distances with centering and scaling:

$$I(s_1 \dots s_n) = \sum_{k < l}^n \frac{I_{s_k, s_l} - \bar{I}_{s_1 \dots s_N, s_1 \dots s_N}}{\sigma(I_{s_1 \dots s_N, s_1 \dots s_N})}$$

Positive values indicate **more independent subsets**; negative values indicate **clustered, redundant models**. This score enables **optimization over model subsets** to maximize independence while maintaining performance and spread (esmvaltool.org).

The **final ClimWIP weight** combines performance and independence multiplicatively:

$$w_i \propto \exp\left(-\frac{d_i^2}{\sigma_d^2}\right) \cdot \left(\sum_j \exp\left(-\frac{S_{ij}^2}{\sigma_S^2}\right)\right)^{-1}$$

where d_i is performance distance, S_{ij} is independence distance, and σ_d, σ_S are tuning parameters controlling weighting strength (esmvaltool.org).

4.3 Independence Detection Methods

4.3.1 Hierarchical Clustering of Model Output Fields Hierarchical clustering of the distance matrix I_{ij} produces **dendrograms revealing model family structures**. Cutting heights indicate similarity thresholds for defining independent clusters. For **CMIP6**, clustering identifies **5-8 major model families**, with effective ensemble sizes of **10-15 rather than nominal 30+** ([GMD](https://esmvaltool.org)). The **cophenetic correlation** between clustering dendrograms and known model relationships validates distance metric choice.

4.3.2 Distance Matrix Analysis for Temperature and Precipitation Analogs Distance matrix visualization as heatmaps exposes **block structures indicating model families and outlier models with unique behavior**. Multidimensional scaling (MDS) or principal component analysis (PCA) of the distance matrix reveals **low-dimensional structure**: often 2-3 dimensions capture

most variance, with the first dimension typically corresponding to **climate sensitivity** and subsequent dimensions capturing **regional pattern differences** ([eLib - DLR electronic library](#)) .

4.3.3 Statistical Tests for Exchangeability Violations Formal tests assess whether the **joint distribution of model predictions is invariant to permutation** of model labels. **Energy statistics** (Székely & Rizzo, 2013) and **kernel-based two-sample tests** detect subtle dependence structures. **Cross-validation partitioning**: if models in the same cluster perform similarly on held-out data, exchangeability is violated ([eLib - DLR electronic library](#)) .

4.4 Application to Interatomic Potential Ensembles

4.4.1 Distance Metrics in Materials Property Space Transferring ClimWIP to materials science requires defining **appropriate distance metrics**. For elastic constant predictions:

$$D_{P_i, P_j}^{\text{prop}} = \sqrt{\frac{1}{M} \sum_{m=1}^M w_m \left(\frac{C_{ij}^{(m)} - \bar{C}^{(m)}}{\sigma_C^{(m)}} \right)^2}$$

where $C_{ij}^{(m)}$ is the elastic constant prediction for material m , with normalization by ensemble mean $\bar{C}^{(m)}$ and variance $\sigma_C^{(m)}$ ([eLib - DLR electronic library](#)) .

Descriptor-space distances capture functional form similarities: **SOAP (Smooth Overlap of Atomic Positions)** fingerprints or **bispectrum coefficients** compared directly between potentials reveal structural similarities even when property predictions differ.

4.4.2 Detecting Shared Functional Forms or Training Data Heritage **Metadata examination** complements statistical detection: potentials derived from the same DFT database, fit with similar loss functions, or implementing the same architectural choices should be flagged for potential dependence. **Statistical detection of near-identical predictions** across material classes—particularly for systems unlikely to be accurately described by simple physics—flags potential **data contamination** from overlapping training sets ([eLib - DLR electronic library](#)) .

4.4.3 Weight Adjustment for Redundant Potential Families If **15 of 23 GLIM potentials** are neural network potentials with shared training data, ClimWIP-style weighting would:

- **Reduce individual NN potential weights** based on mutual similarity
- **Increase weights for structurally distinct potentials** (EAM, MEAM, SNAP)
- **Optimize the performance-independence tradeoff** via cross-validated tuning of σ_d and σ_S

The result: **more balanced, less overconfident ensemble predictions** that better represent true structural uncertainty ([esmvaltool.org](#)) .

5. Ensemble Model Output Statistics (EMOS) and Nonhomogeneous Gaussian Regression (NGR)

5.1 Core Methodology

5.1.1 Post-Processing Raw Ensemble Output to Calibrated Distributions **Ensemble Model Output Statistics (EMOS)** provides a **parametric post-processing framework** that transforms

raw ensemble output into **calibrated probabilistic forecasts**. Unlike BMA, which treats each member as a distinct model requiring individual calibration, EMOS processes the **ensemble as a collective**, extracting summary statistics (mean, variance, spread) that drive predictive distribution parameters. This approach is **computationally efficient and statistically robust**, making it widely deployed in operational weather forecasting (iscram.org) .

5.1.2 NGR as Parametric EMOS for Continuous Variables **Nonhomogeneous Gaussian Regression (NGR)**, developed by **Gneiting et al. (2005)**, is the canonical EMOS formulation for **continuous variables** like temperature, pressure, and elastic constants. The “nonhomogeneous” designation indicates that **predictive distribution parameters vary with the ensemble state**: the location parameter depends on the ensemble mean, while the scale parameter depends on the ensemble spread. This **state-dependent calibration** allows the model to adapt its uncertainty quantification to forecast difficulty (iscram.org) .

5.1.3 Linking Ensemble Moments to Predictive Distribution Parameters The fundamental insight of EMOS/NGR is that **raw ensemble spread is typically a biased and noisy estimator of true predictive uncertainty**. Systematic biases arise from **ensemble undersampling, model errors, and initialization procedures**; random noise arises from **finite ensemble sizes**. EMOS learns a **statistical mapping from ensemble statistics to predictive distribution parameters** that corrects these biases, using training data to estimate optimal mapping coefficients (iscram.org) .

5.2 Mathematical Formulation

5.2.1 Location Parameter: $\mu = a + b \bar{X}$ For exchangeable ensemble members, the NGR location parameter simplifies to:

$$\mu = a + b \cdot \bar{X}$$

where $\bar{X} = \frac{1}{m} \sum_{i=1}^m X_i$ is the **ensemble mean**. The **intercept** a and **slope** b correct systematic bias in the ensemble mean: $b \neq 1$ indicates that the ensemble mean is not optimally calibrated, while $a \neq 0$ indicates additive bias (iscram.org) .

For **non-exchangeable members** (e.g., distinct model families), member-specific coefficients allow differential weighting:

$$\mu = a + b_1 X_1 + \dots + b_m X_m$$

5.2.2 Scale Parameter: $\sigma^2 = c + d S^2$ (**Ensemble Variance**) The **heteroscedastic scale parameter** is:

$$\sigma^2 = c + d \cdot S^2$$

where $S^2 = \frac{1}{m-1} \sum_{i=1}^m (X_i - \bar{X})^2$ is the **ensemble variance**. The **nugget term** $c > 0$ ensures positive variance even when ensemble spread vanishes, representing **irreducible uncertainty**. The **spread coefficient** d is critically important:

d value	Interpretation	Typical situation
$d < 1$	Ensemble spread overestimates uncertainty	Well-designed, overdispersed ensembles (rare)
$d = 1$	Ensemble spread correctly scaled	Ideal case
$d > 1$	Ensemble spread underestimates uncertainty	Typical case: underdispersed raw ensembles

Table 4: Interpretation of NGR spread coefficient d .

In climate applications, $d \approx 1.5 - 3.0$ is **typical**, indicating substantial variance inflation is required for calibration ([iscram.org](https://www.iscram.org)) .

5.2.3 Truncated Normal Variant for Non-Negative Variables For **non-negative variables** (elastic moduli, bulk modulus, shear modulus), the **truncated normal NGR** ensures physical constraints:

$$Y|\mathbf{X} \sim \text{TN}_{[0,\infty)}(\mu, \sigma^2)$$

The truncation ensures **positive predictions** while maintaining computational tractability. Parameter estimation requires modified likelihoods accounting for truncation, typically via numerical optimization or EM-style algorithms ([iscram.org](https://www.iscram.org)) .

5.3 Estimation and Verification

5.3.1 CRPS Minimization for Parameter Estimation **Continuous Ranked Probability Score (CRPS) minimization** is the preferred estimation criterion for EMOS/NGR. The CRPS measures the **integrated squared difference between predictive and empirical CDFs**:

$$\text{CRPS}(F, y) = \int_{-\infty}^{\infty} [F(x) - \mathbb{1}_{[y,\infty)}(x)]^2 dx$$

For **Gaussian predictive distributions**, CRPS has a **closed form**:

$$\text{CRPS}(\mathcal{N}(\mu, \sigma^2), y) = \sigma \left[z(2\Phi(z) - 1) + 2\phi(z) - \frac{1}{\sqrt{\pi}} \right]$$

where $z = (y - \mu)/\sigma$ and Φ, ϕ are standard normal CDF and PDF. This enables **efficient gradient-based optimization** ([iscram.org](https://www.iscram.org)) .

5.3.2 Maximum Likelihood Approaches **Maximum likelihood estimation** provides an alternative, with computational advantages for **complex distributions** (truncated normal, gamma, GEV). For Gaussian NGR:

$$\ell(\theta) = -\frac{1}{2} \sum_{i=1}^n \left[\log(2\pi\sigma_i^2) + \frac{(y_i - \mu_i)^2}{\sigma_i^2} \right]$$

Numerical optimization with constraints ($c \geq 0, d \geq 0$) yields ML estimates. CRPS and ML are **asymptotically equivalent** for large samples; CRPS is generally preferred for its **robustness to distributional assumptions** ([iscram.org](#)) .

5.3.3 Verification Against Competing Post-Processing Methods

Method	Computational cost	Flexibility	Calibration	Typical use case
Raw ensemble	Minimal	None	Poor	Baseline only
BMA	Moderate	High	Excellent	Small ensembles, member-specific info
EMOS/NGR	Low	Moderate	Excellent	Large ensembles, operational use
Quantile regression	Moderate	High	Good	Specific quantile targets
Machine learning	High	Very high	Good (with care)	Abundant training data

Table 5: Comparison of ensemble post-processing methods.

EMOS/NGR consistently achieves **competitive accuracy with substantially lower computational cost** than BMA or machine learning alternatives, making it preferred for **operational high-throughput applications** ([iscram.org](#)) .

5.4 Extension to Materials Property Prediction

5.4.1 Mapping Ensemble Potential Outputs to Elastic Constant Distributions The **direct transfer** of EMOS/NGR to interatomic potential ensembles treats:

- **Ensemble members:** Predictions from different potentials for the same material
- **Ensemble statistics:** Mean and variance across potential predictions
- **Observation:** Reference elastic constant (DFT or experiment)

The NGR equations **learn the mapping** from potential ensemble statistics to calibrated predictive distributions for true elastic constants ([iscram.org](#)) .

5.4.2 Handling Heteroscedastic Errors Across Material Classes **Stratified NGR** with **separate parameter sets** for metals, oxides, semiconductors, etc., captures systematic variation in prediction accuracy. Alternatively, **feature-dependent parameters** where c and d vary with material descriptors (coordination number, band gap, formation energy) enable **continuous adaptation** to predicted difficulty ([AMetSoc](#)) .

5.4.3 Computational Efficiency for High-Throughput Screening Once NGR parameters are estimated, prediction for new materials requires only:

1. Evaluate all potentials (already required for screening)
2. Compute ensemble mean $\bar{X}$ and variance S^2
3. Apply NGR equations: $\mu = a + b\bar{X}$, $\sigma^2 = c + dS^2$

Total overhead: negligible compared to potential evaluation. This enables **real-time calibrated uncertainty quantification** for millions of materials ([iscram.org](#)) .

6. Rank Histograms for Ensemble Calibration Verification

6.1 Diagnostic Purpose and Interpretation

6.1.1 Assessing Uniformity of Observation Ranks Within Ensemble Rank histograms (Tagrand diagrams) provide a **fundamental diagnostic for ensemble calibration** that complements reliability diagrams. The core question: **Is the verifying observation equally likely to fall anywhere within the ensemble range?** For a well-calibrated ensemble, the answer is yes, producing a **uniform (flat) histogram** of observation ranks ([CentAUR](#)) .

6.1.2 Flat Histogram as Indicator of Reliable Spread Uniform rank distribution indicates that **ensemble spread appropriately represents predictive uncertainty**. The observation is as likely to be the smallest, largest, or anywhere between—neither systematically outside the ensemble (indicating underdispersion) nor systematically in the middle (indicating overdispersion) ([CentAUR](#)) .

6.1.3 U-Shaped (Underdispersion) and Dome-Shaped (Overdispersion) Patterns

Histogram Shape	Pattern	Interpretation	Corrective Action
Flat	Uniform	Well-calibrated spread	None needed
U-shaped	Excess at extremes	Underdispersion: ensemble too narrow	Variance inflation ($d > 1$ in NGR)
Dome-shaped	Excess in middle	Overdispersion: ensemble too wide	Variance deflation ($d < 1$)
Skewed	Asymmetric excess	Systematic bias	Location correction ($a \neq 0, b \neq 1$)

Table 6: Rank histogram patterns and their interpretation.

6.2 Mathematical Definition

6.2.1 Rank Computation: Position of Verifying Observation Among Sorted Members For ensemble members $X_1, ..., X_m$ and **observation** Y , the **rank** is:

$$R = 1 + \sum_{i=1}^m \mathbb{1}_{[X_i < Y]}$$

with values in $\{1, 2, ..., m + 1\}$. Ties are handled by **randomization** or **mid-ranks** ([CentAUR](#)) .

6.2.2 Histogram Construction Across Multiple Forecast-Observation Pairs Across N verification instances, **histogram frequencies** are:

$$h_r = \frac{1}{N} \sum_{n=1}^N \mathbb{1}_{[R_n=r]}, \quad r = 1, ..., m + 1$$

Expected frequency under uniformity: $E[h_r] = \frac{1}{m+1}$ ([CentAUR](#)) .

6.2.3 Formal Tests for Uniformity (², Kolmogorov-Smirnov)

- ² **goodness-of-fit**: $\chi^2 = (m + 1)N \sum_{r=1}^{m+1} (h_r - \frac{1}{m+1})^2 \sim \chi_m^2$
- **Kolmogorov-Smirnov**: Maximum deviation of empirical CDF from uniform

Caution: Temporal/spatial correlation reduces effective sample size; **block-bootstrap** or **adjusted critical values** may be needed (CentAUR) .

6.3 Relation to Other Verification Tools

6.3.1 Rank Histograms vs. Reliability Diagrams: Complementary Information

Aspect	Rank Histogram	Reliability Diagram
Assesses	Global distribution calibration	Probability calibration for specific events
Requires	Full ensemble	Probabilistic forecasts (can be derived from ensemble)
Sensitive to	Spread errors, bias	Threshold-specific calibration
Best for	Detecting under/overdispersion	Assessing probability forecast accuracy

Table 7: Complementary roles of rank histograms and reliability diagrams.

6.3.2 Integration with PIT Histograms for Continuous Variables For **parametric predictive distributions** (EMOS/NGR output), the **Probability Integral Transform (PIT)** $U = F(Y) \sim \text{Uniform}(0,1)$ for calibrated forecasts. The **PIT histogram** is the continuous analog of the rank histogram, with similar interpretation (CentAUR) .

6.3.3 Stratified Rank Histograms by Regime or Material Class **Stratification** reveals **conditional calibration**: ensemble reliability may vary by season/region (climate) or by crystal structure/composition (materials). **Stratified histograms** guide **targeted post-processing** or **adaptive ensemble construction** (CentAUR) .

6.4 Implementation for Materials Benchmarks

6.4.1 Defining Ranks When “Truth” Is Uncertain or Multi-Source When **reference values have uncertainty** (multiple DFT functionals, experimental error), **probabilistic ranks** can be computed: the distribution of ranks given reference uncertainty. This yields **rank distributions** rather than point ranks, with summaries (mean, variance, entropy) providing robust calibration assessment (CentAUR) .

6.4.2 Rank Histograms for Cross-Validation in Potential Assessment **Leave-one-material-out cross-validation**:

1. Hold out material \$m\$
2. Train ensemble weights/EMOS parameters on remaining materials
3. Compute rank of held-out reference value in predictive distribution
4. Aggregate ranks across all materials

This yields **out-of-sample calibration assessment** maximizing use of limited benchmark data (CentAUR) .

6.4.3 Detecting Systematic Biases in Specific Potential Families If **EAM potentials consistently underpredict bulk moduli**, observations will have **elevated ranks** (above most EAM predictions). The **rank histogram shows excess at high ranks**, signaling this systematic bias for targeted investigation or exclusion ([CentAUR](#)) .

7. Uncertainty Quantification and Prediction Intervals

7.1 Sources of Uncertainty in Multi-Model Ensembles

7.1.1 Initial Condition Uncertainty **Initial condition uncertainty** arises from imperfect knowledge of system state at prediction initialization. Dominant in **weather forecasting** (chaotic divergence), negligible for **materials property prediction** (equilibrium/ground-state quantities). Analog: **thermal disorder** in finite-temperature properties, treatable via **statistical mechanics ensembles** ([Academia.edu](#)) .

7.1.2 Parameter Uncertainty **Parameter uncertainty** reflects imperfect knowledge of model parameters (empirical coefficients, neural network weights). **Reducible in principle** through additional calibration data, but **practical limits** exist from identifiability and computational cost. Often **implicitly treated** rather than explicitly propagated in interatomic potentials ([Academia.edu](#)) .

7.1.3 Structural Uncertainty: The Dominant Component in Long-Range Projections **Structural uncertainty**—uncertainty about **which model structure is appropriate**—**dominates long-range climate projections and materials property prediction**. No interatomic potential perfectly captures quantum mechanical reality; each embodies **specific approximations with limited domains of validity**. **Cannot be reduced by more data within a fixed model**; requires **model diversity or theoretical advances** ([Academia.edu](#)) .

Uncertainty Source	Climate Science	Materials Science	Reducibility
Initial condition	Dominant (short-range)	Negligible (equilibrium)	Limited (chaos)
Parameter	Moderate	Moderate	Yes (more data)
Structural	Dominant (long-range)	Dominant (novel materials)	No (needs new models)

Table 8: Uncertainty sources and their characteristics in both domains.

7.2 Prediction Interval Construction

7.2.1 Raw Ensemble Spread vs. Calibrated Prediction Intervals **Raw ensemble spread** (range, standard deviation) is **typically overconfident** due to:

- **Underdispersion:** ensemble doesn’t sample all uncertainty sources
- **Model dependence:** correlated members provide less information than nominal size

Calibration factors of 1.5-3× are typically needed for nominal coverage ([Academia.edu](#)) .

7.2.2 BMA-Derived Credible Intervals **BMA predictive distribution** (mixture of member distributions) yields **credible intervals** by numerical integration or simulation:

$$CI_{1-\alpha} = [q_{\alpha/2}, q_{1-\alpha/2}]$$

where q_p is the p -quantile of $p(y|D) = \sum_k w_k p(y|M_k, D)$. **Automatically accounts for within-model and between-model uncertainty** ([Semantic Scholar](#)) .

7.2.3 EMOS/NGR Predictive Distributions and Quantile Extraction NGR predictive distribution $\mathcal{N}(\mu, \sigma^2)$ or $\text{TN}_{[0,\infty)}(\mu, \sigma^2)$ enables **analytic quantiles**:

$$q_p = \mu + \sigma \cdot \Phi^{-1}(p) \quad (\text{Gaussian})$$

or numerical solution for truncated variants. **Computationally efficient** for arbitrary confidence levels ([iscram.org](#)) .

7.3 Challenges in Uncertainty Communication

7.3.1 Overconfidence from Model Dependence Even **calibrated ensembles** may be **overconfident** if **structural diversity is inadequate**. Agreement among correlated models provides **illusory precision**. ClimWIP-style **independence weighting** helps but cannot create missing diversity ([esm-valtool.org](#)) .

7.3.2 Underrepresentation of Structural Uncertainty If **all models share fundamental limitations** (e.g., all local approximations, all trained on same DFT functional), ensemble spread **misses “unknown unknowns”**. **Explicit diversity-seeking** in ensemble construction, including **deliberately unconventional models**, helps ensure coverage ([Academia.edu](#)) .

7.3.3 Temporal Evolution of Uncertainty in Projections **Climate**: Uncertainty typically **increases with lead time** as initial information is lost. **Materials**: Uncertainty **increases with distance from training data** in compositional/structural space. **Adaptive uncertainty quantification** that grows appropriately with “extrapolation distance” is an active research frontier ([Academia.edu](#)) .

7.4 Materials Science Specific Considerations

7.4.1 Propagation from Interatomic Potential to Elastic Constant Uncertainty **Multi-stage uncertainty propagation**:

1. **Potential energy surface uncertainty** (ensemble of potentials)
2. **Equilibrium structure determination** (energy minimization)
3. **Elastic constant calculation** (second derivatives, finite differences)

Each stage **amplifies or correlates errors**; **joint ensemble treatment** preferred to sequential point estimation ([Academia.edu](#)) .

7.4.2 Composition-Structure-Property Uncertainty Chaining **Materials discovery workflow uncertainty**:

- **Composition** → **stable structure(s)?** (structure prediction uncertainty)
- **Structure** → **property values?** (interatomic potential uncertainty)
- **Properties** → **performance in application?** (processing, environment, degradation)

Propagating uncertainty through this chain is essential for rational decision-making ([Academia.edu](#)) .

7.4.3 Decision-Making Under Uncertainty for Materials Discovery Value-of-information analysis: Is additional computation (more accurate DFT, more potentials, more materials) warranted given current uncertainty?

Robust optimization: Identify materials performing well across **range of plausible property values**, not just at point predictions.

Multi-objective balance: Performance, uncertainty, cost, synthesizability ([Academia.edu](#)) .

8. Transferable Mathematical Frameworks and Implementation Roadmap

8.1 Directly Transferable Methods

8.1.1 BMA Weighting with Potential-Specific Posterior Probabilities Direct application to GLIM-scale ensembles:

- **Training data:** DFT benchmark calculations
- **Likelihood:** Gaussian or heavy-tailed error distributions
- **Estimation:** EM algorithm or MCMC
- **Output:** Weighted combinations with calibrated uncertainties ([Semantic Scholar](#))

8.1.2 EMOS/NGR Recalibration for Elastic Constant Predictions Direct application:

- **Inputs:** Ensemble mean $\bar{X}$, variance S^2 across potential predictions
- **Parameters:** (a, b, c, d) estimated by CRPS minimization
- **Output:** Calibrated Gaussian or truncated normal predictive distributions ([iscram.org](#))

8.1.3 Rank Histogram and Reliability Diagram Diagnostics Identical construction and interpretation, with DFT as “observations.” **Stratification by material class** reveals conditional calibration ([CentAUR](#)) .

8.2 Adaptation Requirements

8.2.1 Defining Appropriate “Observations” for Calibration Challenge: No single “ground truth” for most materials properties.

Approaches:

- **Hierarchical Bayesian:** Multiple fidelity levels with explicit error models
- **Consensus estimation:** Combine multiple reference sources
- **Cross-validation:** Pseudo-observations from held-out data ([Academia.edu](#))

8.2.2 Handling Multi-Fidelity Data (DFT, Experiment, Different XC Functionals) Multi-fidelity Gaussian processes or hierarchical BMA:

- **Fidelity as model dimension:** Separate weights for PBE, PBEsol, SCAN, HSE06, experiment
- **Discrepancy modeling:** Explicit bias functions relating lower to higher fidelity ([Academia.edu](#))

8.2.3 Computational Cost-Benefit Analysis for Large Ensembles

Method	Training cost	Prediction cost	Best for
BMA	Moderate (EM iterations)	Moderate (mixture evaluation)	Small ensembles, rich member info
EMOS/NGR	Low (closed-form CRPS)	Negligible (linear transform)	Large-scale screening
Machine learning	High (neural network training)	Low (forward evaluation)	Abundant data, complex patterns

Table 9: Computational trade-offs for ensemble post-processing methods.

8.3 Emerging Opportunities

8.3.1 Machine Learning Integration (Random Forest Ensemble Averaging) Random forests as ensemble method: diverse trees trained on bootstrap samples, predictions averaged. **Natural bridge** between climate-style ensembles and modern ML. **Uncertainty quantification** via out-of-bag estimation or quantile regression forests ([Academia.edu](https://www.academia.edu)) .

8.3.2 Deep Learning for Non-Parametric Post-Processing Neural networks learn complex mappings from ensemble statistics to predictive distributions. **Attention mechanisms** identify which ensemble members to trust for which inputs. **Risk:** Overfitting, loss of interpretability; **mitigation:** Careful validation, regularization ([Academia.edu](https://www.academia.edu)) .

8.3.3 Active Learning for Adaptive Ensemble Construction **Sequentially select** which potentials to evaluate on which materials based on **estimated information gain**. **Acquisition functions** from Bayesian optimization: maximize uncertainty reduction, maximize expected improvement, maximize information gain. **Closes the loop** between ensemble construction and combination ([Academia.edu](https://www.academia.edu)) .

9. Key Literature and References

9.1 Foundational BMA Work

Raftery, A. E., Gneiting, T., Balabdaoui, F., & Polakowski, M. (2005). Using Bayesian model averaging to calibrate forecast ensembles. *Monthly Weather Review*, 133(5), 1155–1174. <https://doi.org/10.1175/MWR2906.1>

Fraley, C., Raftery, A. E., & Gneiting, T. (2010). Calibrating multimodel forecast ensembles with exchangeable and missing members using Bayesian model averaging. *Monthly Weather Review*, 138(5), 1904–1915. <https://doi.org/10.1175/2009MWR3046.1>

9.2 CMIP5/CMIP6 Ensemble Literature

Knutti, R., & Tebaldi, C. (2010). Challenges in combining projections from multiple climate models. *Journal of Climate*, 23(10), 2739–2758. <https://doi.org/10.1175/2009JCLI3361.1>

Knutti, R., Sedláček, J., Sanderson, B. M., Lorenz, R., Fischer, E. M., & Eyring, V. (2017). A climate model projection weighting scheme accounting for performance and interdependence. *Geophysical Research Letters*, 44(4), 1909–1918. <https://doi.org/10.1002/2016GL072012>

Brunner, L., Pendergrass, A. G., Lehner, F., Merrifield, A. L., Lorenz, R., & Knutti, R. (2020). Reduced global warming from CMIP6 projections when weighting models by performance and independence. *Earth System Dynamics*, 11(4), 995–1012. <https://doi.org/10.5194/esd-11-995-2020>

Merrifield, A. L., Brunner, L., Lorenz, R., Sedláček, J., & Knutti, R. (2020). An investigation of weighting schemes suitable for incorporating large ensembles into multi-model ensembles. *Earth System Dynamics*, 11(4), 807–834. <https://doi.org/10.5194/esd-11-807-2020>

Merrifield, A. L., Brunner, L., & co-authors. (2023). Climate model Selection by Independence, Performance, and Spread (ClimSIPS). *Geoscientific Model Development*, 16, 4715–4747. <https://doi.org/10.5194/gmd-16-4715-2023>

9.3 EMOS/NGR and Post-Processing

Gneiting, T., Raftery, A. E., Westveld, A. H., & Goldman, T. (2005). Calibrated probabilistic forecasting using ensemble model output statistics and minimum CRPS estimation. *Monthly Weather Review*, 133(5), 1098–1118. <https://doi.org/10.1175/MWR2904.1>

Thorarinsdottir, T. L., & Gneiting, T. (2010). Probabilistic forecasts of wind speed: Ensemble model output statistics by using heteroscedastic censored regression. *Journal of the Royal Statistical Society: Series A*, 173(2), 371–388. <https://doi.org/10.1111/j.1467-985X.2009.00616.x>

9.4 Verification and Calibration Diagnostics

Hamill, T. M. (2001). Interpretation of rank histograms for verifying ensemble forecasts. *Monthly Weather Review*, 129(3), 550–560. [https://doi.org/10.1175/1520-0493\(2001\)129<0550:IORHFV>2.0.CO;2](https://doi.org/10.1175/1520-0493(2001)129<0550:IORHFV>2.0.CO;2)

Hamill, T. M., & Juras, J. (2006). Measuring forecast skill: Is it real skill or is it the varying climatology? *Quarterly Journal of the Royal Meteorological Society*, 132(621C), 2905–2923. <https://doi.org/10.1256/qj.06.25>

Wilks, D. S. (2019). *Statistical methods in the atmospheric sciences* (4th ed.). Academic Press. <https://doi.org/10.1016/C2017-0-03921-6>

Dimitriadis, T., Gneiting, T., Jordan, A. I., & Vogel, P. (2023). Stable reliability diagrams for probabilistic classifiers. *Proceedings of the National Academy of Sciences*, 120(8), e2016191118. <https://doi.org/10.1073/pnas.2016191118>